

High Performance Machine Learning Lab 3

Name: Vishwajeet Kulkarni
NYU ID: vk2630

Prob 1: Chatbot Seq-2-Seq Model

1. Wandb Project:

https://wandb.ai/vk2630-new-york-university/LAB3_HPML?nw=nwuservk2630

2. Hyperparameter Sweeps:

Strategy: Random Search

Total Runs: 25

Overview

Workspace

Runs

Agents

Controls

Logs

Vk2630's sweep workspace

Personal workspace

Autosaved just now

Sweep: Lab3-sweep

Search runs

Filter

Group

Sort

New sweep

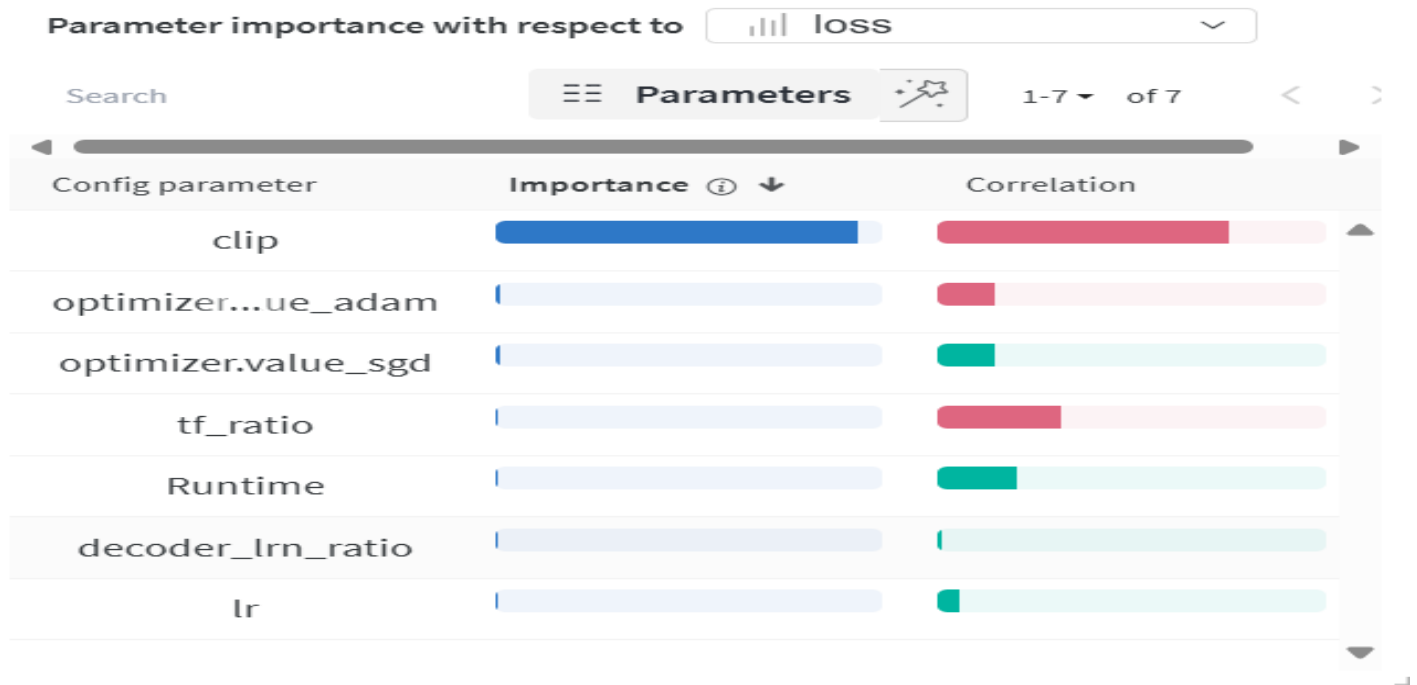
Columns

</


I observed that run with id vk2630 has least loss of . Hyperparameters value of above model:

1. Learning Rate: 0.00025
2. Clip: 100.0
3. Teacher Forcing Ratio: 1
4. Decoder Learning Ratio: 10
5. Optimizer: adam

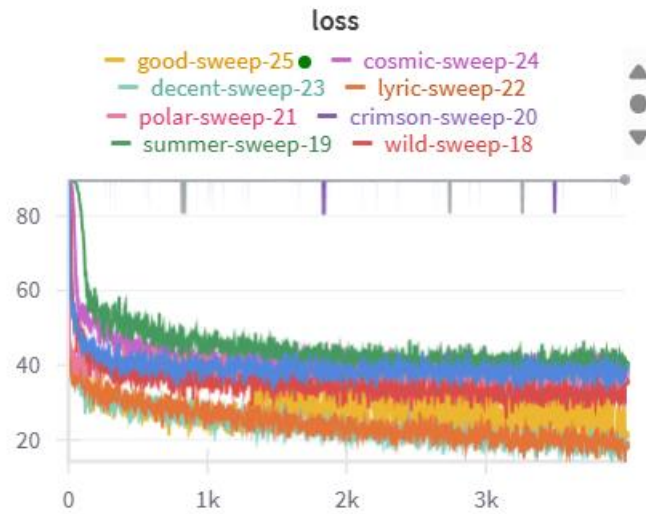
3. Feature Importance:



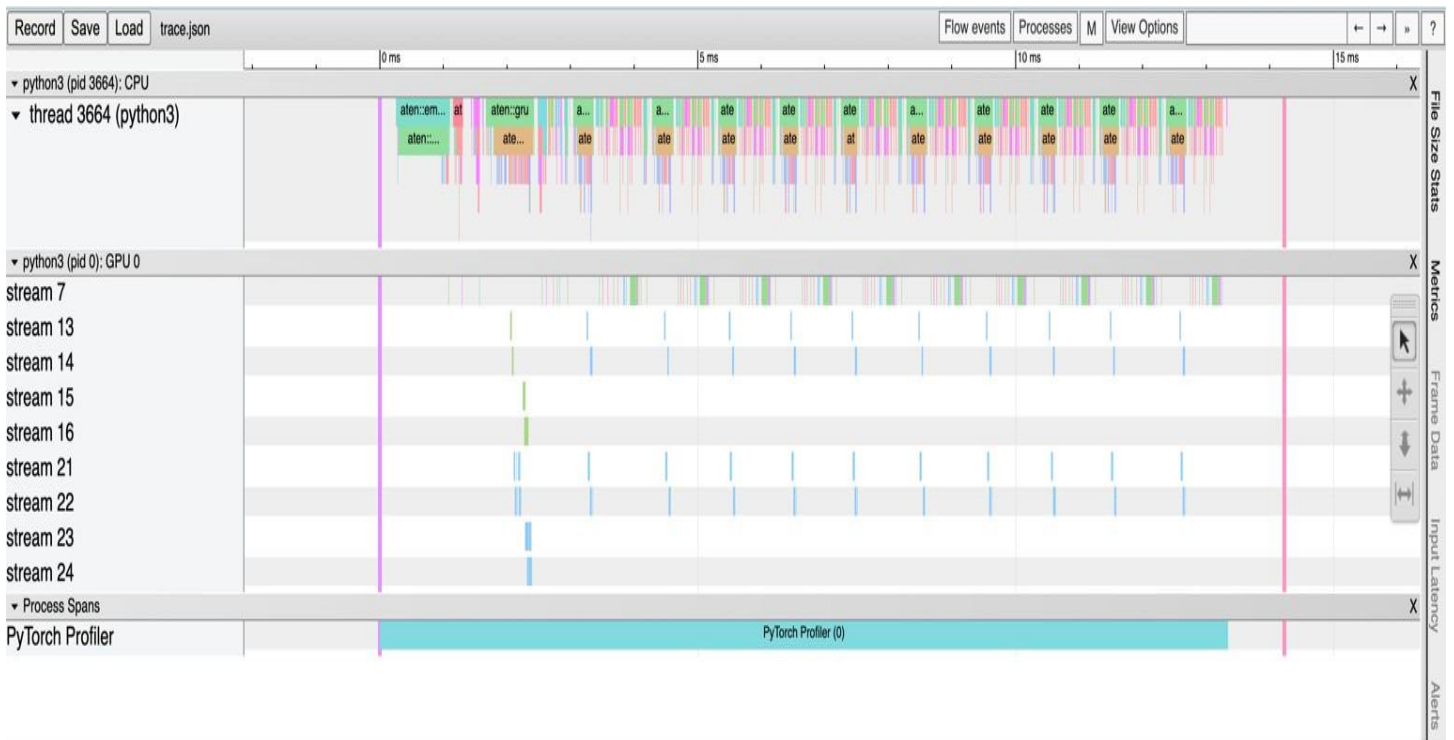
From above panel, we can observe that value of clip has very high effect on the loss i.e if value of clip is high, loss will be less (Value of correlation is negative). Moreover, we can see that both “ADAM” and “SGD” has high correlation, but ADAM has positive correlation and hence ADAM will be better choice.



4. Variation of loss with iterations



5. PyTorch Profiler Tracing



6. Measurement of time and memory Consumption of model's operators

| | Name | Self CPU | CPU total | CPU time avg | CUDA total | CUDA time avg | CPU Mem | Self CPU Mem | CUDA Mem | Self CUDA Mem |
|---------------------------------|-----------------------------|-----------|-----------|--------------|------------|---------------|---------|--------------|-----------|---------------|
| | aten::empty | 523.000us | 523.000us | 6.226us | 0.000us | 0.000us | 24 b | 24 b | 335.39 Mb | 335.39 Mb |
| | aten::embedding | 4.636ms | 6.319ms | 574.455us | 56.000us | 5.091us | 0 b | 0 b | 24.00 Kb | 0 b |
| | aten::reshape | 27.000us | 41.000us | 3.727us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::view | 31.000us | 31.000us | 1.292us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::index_select | 1.028ms | 1.631ms | 148.273us | 56.000us | 5.091us | 0 b | 0 b | 24.00 Kb | 0 b |
| | aten::resize_ | 99.000us | 99.000us | 9.000us | 0.000us | 0.000us | 0 b | 0 b | 24.00 Kb | 24.00 Kb |
| | cudaLaunchKernel | 33.995ms | 33.995ms | 157.384us | 209.000us | 0.968us | 0 b | 0 b | 0 b | 0 b |
| ous namespace)::indexSelectS... | | 0.000us | 0.000us | 0.000us | 56.000us | 5.091us | 0 b | 0 b | 0 b | 0 b |
| | aten::to | 3.000us | 3.000us | 3.000us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::_pack_padded_sequence | 51.000us | 964.000us | 964.000us | 3.000us | 3.000us | 16 b | 0 b | 4.00 Kb | 0 b |
| | aten::slice | 201.000us | 210.000us | 16.154us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::as_strided | 69.000us | 69.000us | 0.476us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::cat | 1.278ms | 1.735ms | 55.968us | 133.000us | 4.290us | 0 b | 0 b | 54.00 Kb | 54.00 Kb |
| | aten::narrow | 7.000us | 16.000us | 16.000us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | cudaMemcpyAsync | 49.000us | 49.000us | 24.500us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| Memcpy DtoD (Device -> Device) | | 0.000us | 0.000us | 0.000us | 6.000us | 3.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::select | 101.000us | 113.000us | 5.136us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::item | 8.000us | 10.000us | 5.000us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::_local_scalar_dense | 4.000us | 4.000us | 2.000us | 0.000us | 0.000us | 0 b | 0 b | 0 b | 0 b |
| | aten::zeros | 27.000us | 1.613ms | 537.667us | 2.000us | 0.667us | 0 b | 0 b | 8.00 Kb | 0 b |

Prob 2: TorchScript Seq-2-Seq Model

Q1. Explain the differences between tracing and scripting and how they are used in TorchScript?

Solution:

In **tracing**, both the model and sample input data are required to record computations and generate a graph-based function. However, it only logs operations specific to the provided input, making it ineffective for handling data-dependent control flow. In **scripting**, the model alone is converted into TorchScript without requiring sample data. This process translates the model code into a subset of Python that retains all control flow structures. For models without control flow dependencies, `torch.jit.trace()` can be used directly without modifications, as it automatically converts the model into TorchScript. However, scripting may require adjustments to the model code to ensure compatibility with TorchScript syntax before applying `torch.jit.script()`. When using tracing, the model's device and dropout layers must be set to test mode beforehand, as the traced model does not manage these operations internally. In contrast, scripting allows these settings to be adjusted just before inference, similar to standard eager execution.

Q2. Explain the changes needed in the chatbot model to allow for scripting.

Solution:

In the given model, there are three sub-modules: Encoder, Decoder, and *GreedySearchDecoder*.

Changes required:

1. **Explicit Type Annotations for Forward Method Arguments:** By default, TorchScript assume all parameters of function as tensor. Hence, in case we need to pass any argument with different type like int in our case, we need to specify the type in python function.
2. **Passing decoder_n_layers as a Constructor Argument:** Earlier, it was using fetching this value from decoder but since we are using traced version of decoder we will not be able to access that value anymore and hence we need to pass this to its constructor for it to use.

```
#Modified GreedySearchDecoder for scripting the module  
  
class GreedySearchDecoderScript(torch.jit.ScriptModule):  
    def __init__(self, encoder, decoder, decoder_n_layers):
```

Q3: Comparing Latency:

| | Latency on CPU (ms) | Latency on GPU (ms) |
|-------------|---------------------|---------------------|
| Pytorch | 345.456431 | 10.838052 |
| Torchscript | 13.652194 | 19.023610 |
| SpeedUp | 25.304097 | 0.569716 |
